


Task Difficulty Distribution by Benchmark



Difficulty (b) estimated via 2-parameter logistic IRT (py-irt, hierarchical priors, 1000 epochs SVI).
Red line = median. Tasks never solved by any model have prior-regularized b estimates.
Data: Liu et al., "BRIDGE: Predicting Human Task Completion Time From Model Performance" (arXiv:2602.07267, 2026).