


MLE-bench: CDF of IRT Difficulty



Difficulty (b) estimated via 2-parameter logistic IRT (py-irt, hierarchical priors, 1000 epochs SVI).
114 MLE-bench items scored binary pass/fail across 176 model+scaffold submissions.
61 items never solved by any model; their b values are prior-regularized extrapolations.