


SWE-bench: CDF of IRT Difficulty



*Difficulty (b) estimated via 2-parameter logistic IRT (py-irt, hierarchical priors, 1000 epochs SVI).
500 SWE-bench items scored binary pass/fail across 176 model+scaffold submissions.
32 items never solved by any model; their b values are prior-regularized extrapolations.*