


SWE-bench (single-benchmark IRT): CDF of IRT Difficulty



Difficulty (b) estimated via 2-parameter logistic IRT (py-irt, hierarchical priors, 1000 epochs SVI).
500 SWE-bench (single-benchmark IRT) items scored binary pass/fail across 176 model+scaffold submissions.
32 items never solved by any model; their b values are prior-regularized extrapolations.
Data: Liu et al., "BRIDGE: Predicting Human Task Completion Time From Model Performance" (arXiv:2602.07267, 2026).